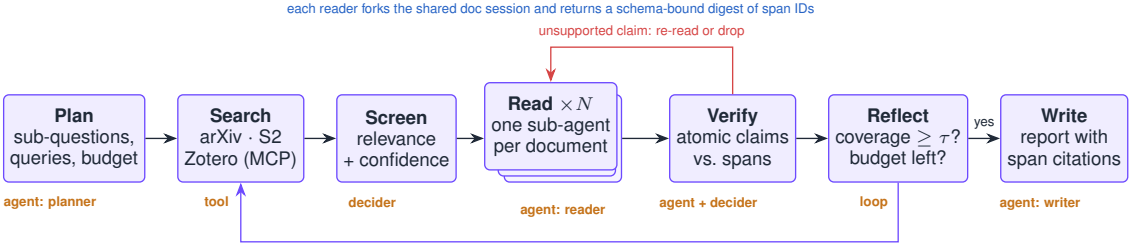
each reader forks the shared doc session and returns a schema-bound digest of span IDs
unsupported claim: re-read or drop
Plan
sub-questions, queries, budget
Search
arXiv · S2
Zotero (MCP)
Screen
relevance + confidence
Read $\times N$
one sub-agent per document
Verify
atomic claims vs. spans
Reflect
coverage $\geq \tau$?
budget left?
yes
Write
report with span citations
agent: planner
tool
decider
agent: reader
agent + decider
loop
agent: writer
no: new queries / citation chasing / read more